

Simulating Virtual Output Queues

Speaker: Wenkai Dai (CT Group)

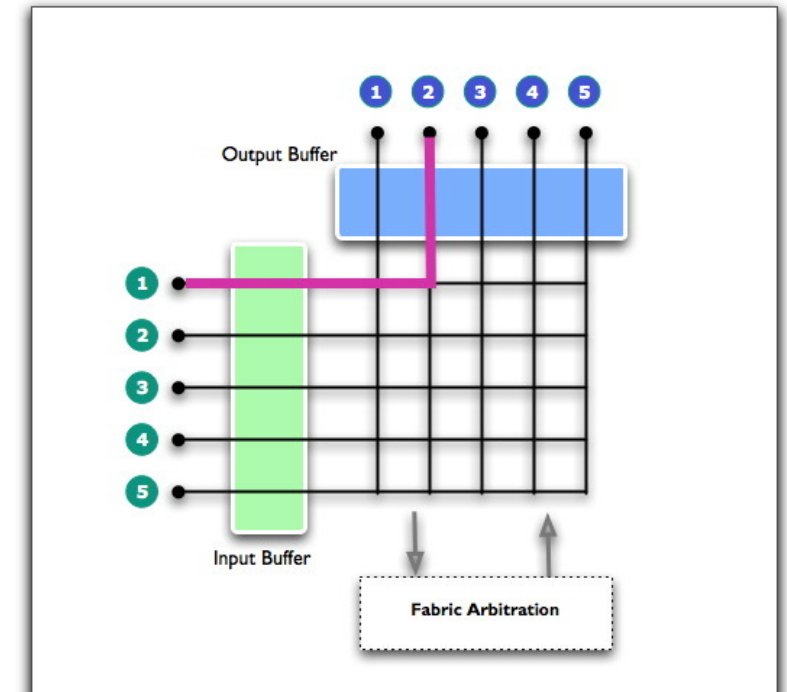




Motivation

-



Packet Switch Network



Switches Forward Packets

- **Infrastructure of Switch:**
 - input links: (I_1, \dots, I_n)
 - output links: (O_1, \dots, O_n)
 - Switch fabric: copying packets from input links to output links
- **Header of Packet** p : $(I[p], O[p])$ indicates input and output links respectively
- **Forwarding Steps:**
 1. a packet $p = (I[p], O[p])$ arrives at an input link $I[p]$
 2. moving p to the output link $O[p]$
 3. p leaves the output link $O[p]$
- Each input/output link can only have one packet processed at a time step_____

Model 1: Virtual Output Queueing (VOQ)


- Each output link contains a buffering queue
- No time cost for switching fabric
- Only one packet can leave at each output link at each time step

One step under pure output queueing:

Packets arrive on input links

Each packet p of type $(I[p], O[p])$ is moved to output buffer $O[p]$

At most one packet departs from each output buffer



Model 2: Input/Output Queueing

- Each input/output link contains a queue
- Switching fabric need a time step to move packets
- Only one packet can leave at each output link at each time step


One step under input/output queueing:

Packets arrive on input links and are placed in input buffers


A set of packets whose types form a matching are moved to their associated output buffers

At most one packet departs from each output buffer



It's a  Match!



Head-of-Line Blocking



Epilogue: Algorithms that run forever



Simulate Virtual Queueing By Input/Output Queueing

- Speed-up switching fabric by factor of two is enough

One step under sped-up input/output queueing:

Packets arrive on input links and are placed in input buffers

A set of packets whose types form a matching are moved to their associated output buffers

At most one packet departs from each output buffer

A set of packets whose types form a matching are moved to their associated output buffers

Preliminaries of Proof

- $TL(p)$: the leave time of p under VOQ model (deadline)
- Output queues: unordered and a packet can leave at any place
- Input queues: ordered by deadlines, but arbitrarily insert and move to the head
- Input cushion $IC(p)$: the number of packets before p in its input queue
- Output cushion $OC(p)$: packets with earlier deadline in the output queue of p
- Define $Slack(p) = OC(p) - IC(p)$ (indicate the slackness, and larger is better)

Invariants To Be Guaranteed

- (i) $Slack(p) \geq 0$ for all unprocessed packets p .
- (ii) In any step that begins with $IC(p) = OC(p) = 0$, packet p will be moved to its output buffer in the first matching.


(E.1) *If properties (i) and (ii) are maintained for all unprocessed packets at all times, then every packet p will depart at its time to leave $TL(p)$.*

Proof. If p is in its output buffer at the start of step $TL(p)$, then it can clearly depart. Otherwise it must be in its input buffer. In this case, we have $OC(p) = 0$ at the start of the step. By property (i), we have $Slack(p) = OC(p) - IC(p) \geq 0$, and hence $IC(p) = 0$. It then follows from property (ii) that p will be moved to the output buffer in the first matching of this step, and hence will depart in this step as well. ■

Initialization of Invariant

Moving a Matching through a Switch When a packet p first arrives on an input link, we insert it as far back in the input buffer as possible (potentially somewhere in the middle) consistent with the requirement $Slack(p) \geq 0$. This makes sure property (i) is satisfied initially for p .

Next to Show $Slack(p) = OC(p) - IC(p)$ Increasing



- Note only when $OC(p)$ (decrease) or $IC(p)$ (increase) then $Slack(p)$ decrease
- A packet arrive/leave might decrease $Slack(p)$ by one (in total two)
- Try to increase $Slack(p)$ by one for moving each matching
- Select stable matchings
 - Preference of output O: the input I where the packet has the earliest deadline
 - Preference of input I: the packet to O is the forwardmost than other output O'
 - Stability: no other matching can satisfy the preference better

Moving A Matchings Ensures $Slack(p)$ Increasing by 1

(E.2) *Suppose the switch always moves a stable matching M with respect to the preference lists defined above. (And for each type (I, O) contained in M , we select the packet of this type with the earliest time to leave). Then, for all unprocessed packets p , the value $Slack(p)$ increases by at least 1 when the matching M is moved.*

Proof. Consider any unprocessed packet p . Following the discussion above, suppose that no packet ahead of p in $I[p]$ is moved as part of the matching M , and no packet destined for $O[p]$ with an earlier time to leave is moved as part of M . So, in particular, the pair $(I[p], O[p])$ is not in M ; suppose that pairs $(I', O[p])$ and $(I[p], O')$ belong to M .

Now p has an earlier time to leave than any packet of type $(I', O[p])$, and it comes ahead of every packet of type $(I[p], O')$ in the ordering of $I[p]$. It follows that $I[p]$ prefers $O[p]$ to O' , and $O[p]$ prefers $I[p]$ to I' . Hence the pair $(I[p], O[p])$ forms an instability, which contradicts our assumption that M is stable. ■